

Università degli
Studi di Milano-Bicocca




Convolutional Neural Networks

Prof. Flavio Piccoli - Dott. Mirko Paolo Barbato

So far..

..we used datasets with (mostly) unrelated variables





But in real world data exist many types of data that have neighborhood dependency

Time-Series

- Time-series: same variable changing in the time



Example: the price of a stock



Signals

- Same quantity acquired in different frequencies, spatial locations, etc.


Example: the hyperspectral signal




Can we use something smarter than NNs?

- Yes!
- We can use Convolutional Neural Networks
- They exploit local dependencies

Linear Layer




Convolutional Layer




Can we use something smarter than NNs?

- Yes!
- We can use Convolutional Neural Networks
- They exploit local dependencies

Linear Layer




Convolutional Layer




Can we use something smarter than NNs?

- Yes!
- We can use Convolutional Neural Networks
- They exploit local dependencies

Linear Layer




Convolutional Layer




Can we use something smarter than NNs?

- Yes!
- We can use Convolutional Neural Networks
- They exploit local dependencies

Linear Layer




Convolutional Layer



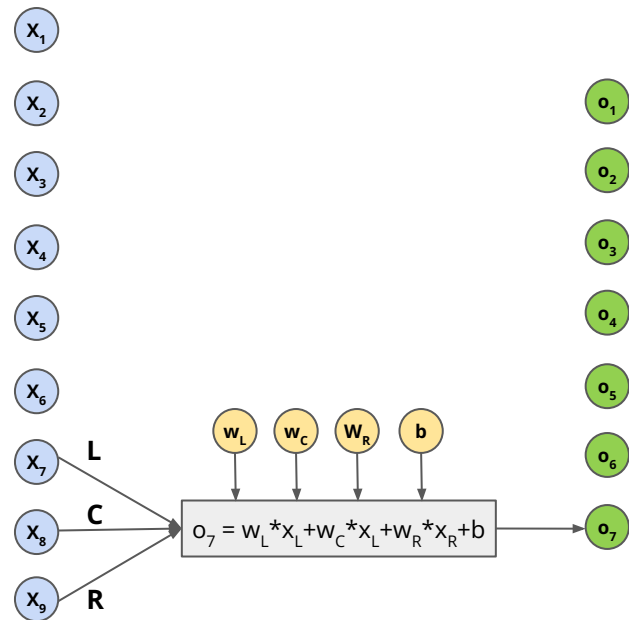
Can we use something smarter than NNs?

- Yes!
- We can use Convolutional Neural Networks
- They exploit local dependencies

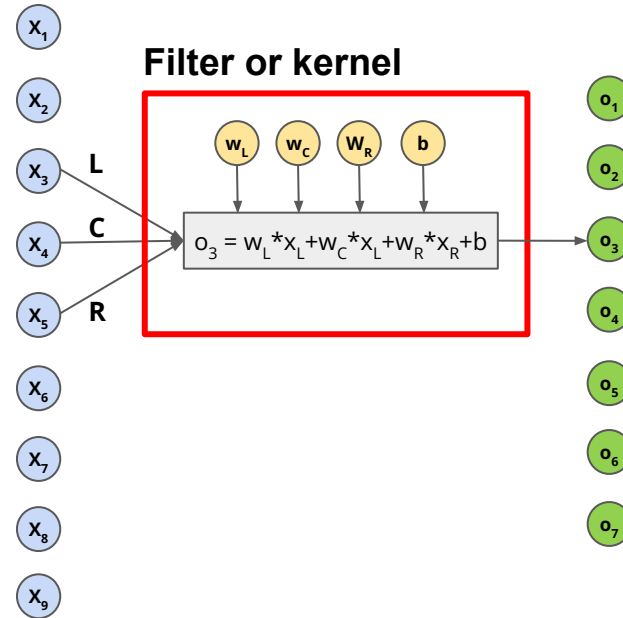
Linear Layer



Convolutional Layer



Convolutional layer - ConvLayer1 D



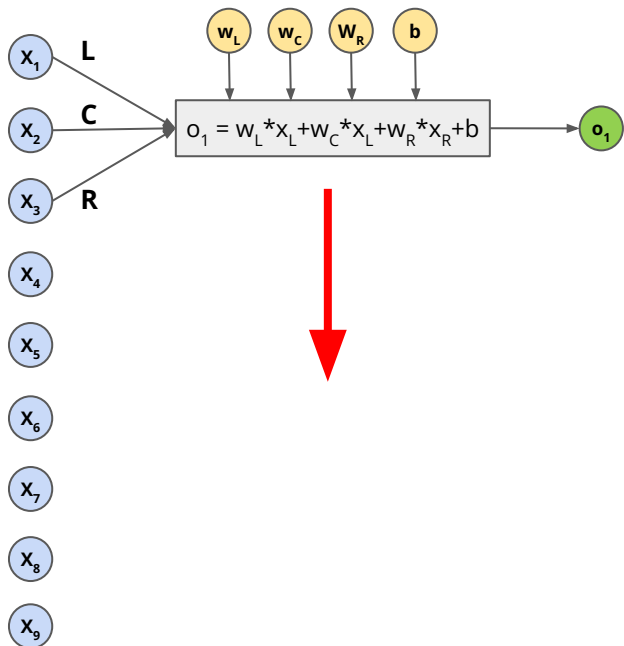
kernel size = number of neighbours considered at once. Typically 3.

stride = step of application. Usually 1 or 2.


Stride

- it's the length of the step by which the filter moves

Stride 2




Stride 1




Stride

Stride 2




Stride 1




Stride

Stride 2



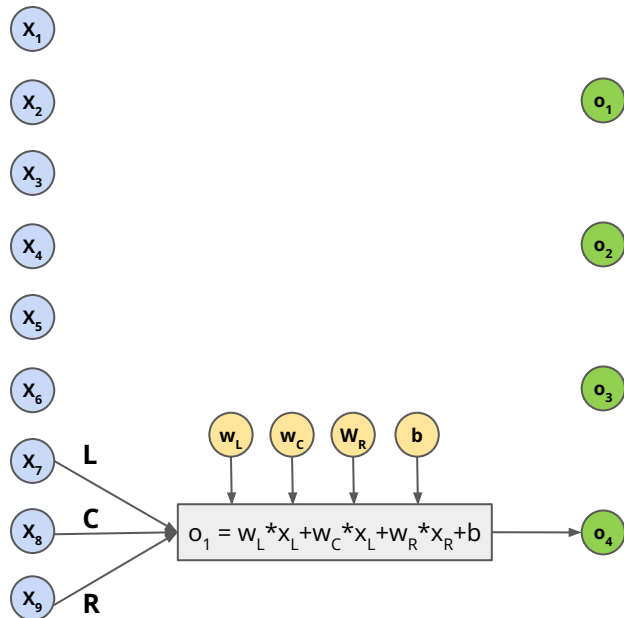
Stride 1



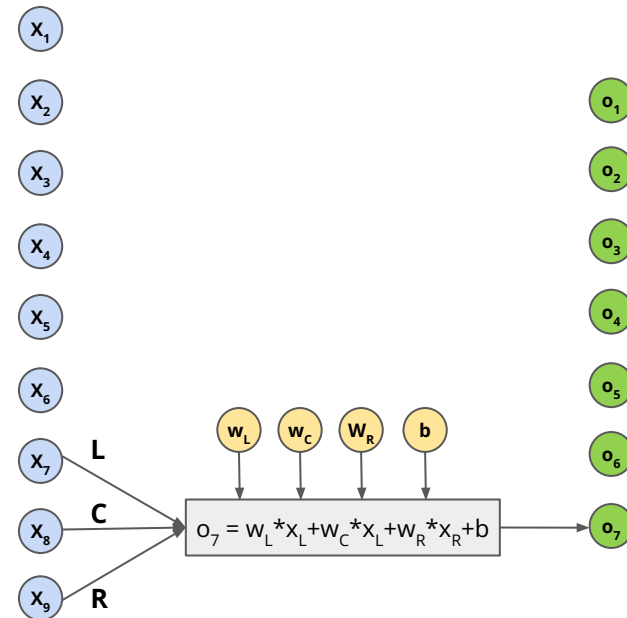
Stride

- stride 2 halves the size

Stride 2



Stride 1

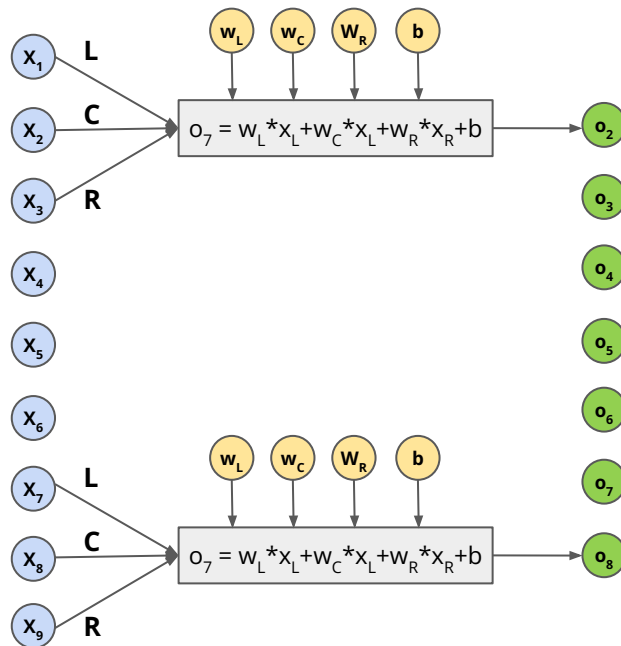


Padding

- Suppose you want an output having the same size of the input
- You need to add padding

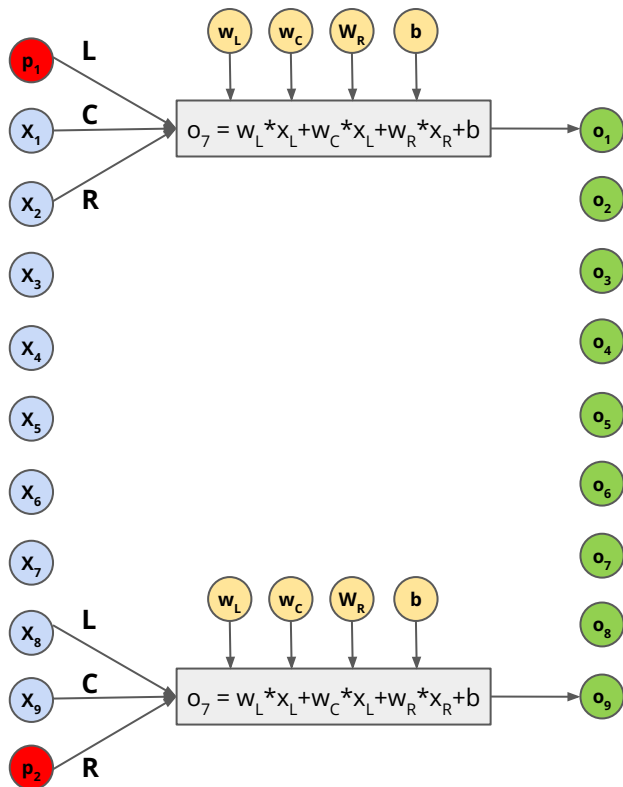
9 elements

7 elements



Padding

- Suppose you want an output having the same size of the input
- You need to add padding



9 elements
+ 2 pad.

9 elements

There are several types of padding:

- constant
- reflect
- replicate
- circular


Default is 'constant' with value 0

To maintain the same dimension, pad should be:

$$pad = \frac{kernel\ size - 1}{2}$$

Convolutional layer - ConvLayer1D

- Of course we can use multiple filters at once and obtain several activations



`torch.nn.Conv1d(in_channels, out_channels, kernel_size, stride=1, padding=1)`

Number of input channels

Input usually has 1.


Number of filters

Determines the
number of outputs

Structure of a CNN


- Suppose to have a batch of 2 samples.
 - Each sample is a signal of 9 elements (1ch)
- Suppose that convolutional layers produce activations of 5 elements having 8 channels

(2,1,9)



Structure of a CNN

- Feature extraction uses convolutional layers to extract meaningful features
- Classification uses fully connected layers (linear layers) to take a decision



Example in Pytorch

```
class CNN(nn.Module):
```

```
    def __init__(self):
```

```
        # initialize super class
```

```
        super(CNN, self).__init__()
```

```
        # define conv layers
```

```
        self.layer1 = nn.Conv1d( 1, 32, kernel_size=3, stride=2, padding=1)
```

```
        self.layer2 = nn.ReLU()
```

```
        self.layer3 = nn.Conv1d(32, 64, kernel_size=3, stride=2, padding=1)
```

```
        self.layer4 = nn.ReLU()
```

```
        self.layer5 = nn.Conv1d(64, 128, kernel_size=3, stride=2, padding=1)
```

```
        self.layer6 = nn.ReLU()
```

```
        # define linear layer
```

```
        self.layer7 = nn.Linear(384, 1)
```

```
    def forward(self, x):
```

```
        # apply convolution layers
```

```
        x = self.layer1(x)
```

```
        x = self.layer2(x)
```

```
        x = self.layer3(x)
```

```
        x = self.layer4(x)
```

```
        x = self.layer5(x)
```

```
        x = self.layer6(x)
```

```
        # reshape from (10, 128, 3) to (10, 384)
```

```
        x = x.reshape(x.shape[0], -1)
```

```
        # fully connected
```

```
        x = self.layer7(x)
```

```
        # return output
```

```
        return x
```

```
# create cnn
```

```
cnn = CNN()
```

```
# create fake input
```

```
inp = torch.rand(10, 1, 20)
```

```
# compute output
```

```
out = cnn(inp)
```

Example in Pytorch

```
class CNN(nn.Module):
```

```
    def __init__(self):
```

```
        # initialize super class
```

```
        super(CNN, self).__init__()
```

```
        # define conv layers
```

```
        self.layer1 = nn.Conv1d( 1, 32, kernel_size=3, stride=2, padding=1)
```

```
        self.layer2 = nn.ReLU()
```

```
        self.layer3 = nn.Conv1d(32, 64, kernel_size=3, stride=2, padding=1)
```

```
        self.layer4 = nn.ReLU()
```

```
        self.layer5 = nn.Conv1d(64, 128, kernel_size=3, stride=2, padding=1)
```

```
        self.layer6 = nn.ReLU()
```

```
        # define linear layer
```

```
        self.layer7 = nn.Linear(384, 1)
```

```
    def forward(self, x):
```

```
        # apply convolution layers
```

```
        x = self.layer1(x)
```

```
        x = self.layer2(x)
```

```
        x = self.layer3(x)
```

```
        x = self.layer4(x)
```

```
        x = self.layer5(x)
```

```
        x = self.layer6(x)
```

```
        # reshape from (10, 128, 3) to (10, 384)
```

```
        x = x.reshape(x.shape[0], -1)
```

```
        # fully connected
```

```
        x = self.layer7(x)
```

```
        # return output
```

```
        return x
```

```
# create cnn
```

```
cnn = CNN()
```

```
# create fake input
```

```
inp = torch.rand(10, 1, 20)
```

```
# compute output
```

```
out = cnn(inp)
```

Convolutional layers

Last activation has shape (10, 128, 3)

Reshaping

to shape (10, 384). 384 is 128*3.

Linear layers

Project 128 elements in 1c, the output

Batch normalization

This layer controls the range of the activations

A layer with weights to be learned, is easier (or sometime it doesn't work otherwise) if input data has:

- mean = 0
- stddev = 1

$$bnorm(x) = \frac{x - \text{mean}(x)}{\text{stddev}(x)}$$

Running mean

Running standard deviation

Batch normalization layer is used many times inside a network

In PyTorch:


```
torch.nn.BatchNorm1d(num_features, eps=1e-05, momentum=0.1)
```

Layer order

Which is the correct order of the conv, ReLU, BatchNorm layers?



OR




Layer order

Which is the correct order of the conv, ReLU, BatchNorm layers?



OR



There is not a better order, depends from the task / dataset

Layer order


Which is the correct order of the conv, ReLU, BatchNorm layers?



In this configuration, at the end activations are normalized

Layer order

Which is the correct order of the conv, ReLU, BatchNorm layers?



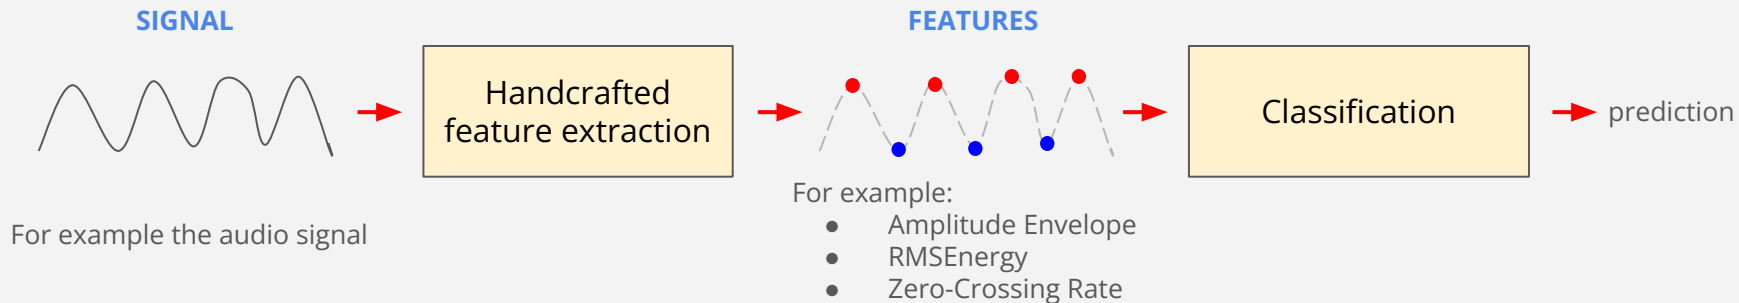
In this configuration instead, at the end activations are sparse (i.e. full of zeros)

Old approaches vs convolutional neural networks

The main difference is that in traditional approaches you have to select the most appropriate features for the task while with CNNs this is done automatically

TRADITIONAL APPROACH


REQUIRES DOMAIN KNOWLEDGE FEATURE ENGINEERING



DEEP-LEARNING-BASED APPROACH

DOES NOT REQUIRE ANY DOMAIN KNOWLEDGE

AUTOMATIC FEATURE EXTRACTION + CLASSIFICATION





The hyperspectral signal

Reflectance vs Irradiance

Reflectance

- is the measure of the proportion of electromagnetic energy reflected by a surface material
- is unitless and measured in percentages

Irradiance

- is the amount of energy received by a surface per unit area
- Cameras measure irradiance

Surface reflectance




Image courtesy by Merav Mazouz & Matan Kolath

Reflectance model

We are interested in the reflected radiation

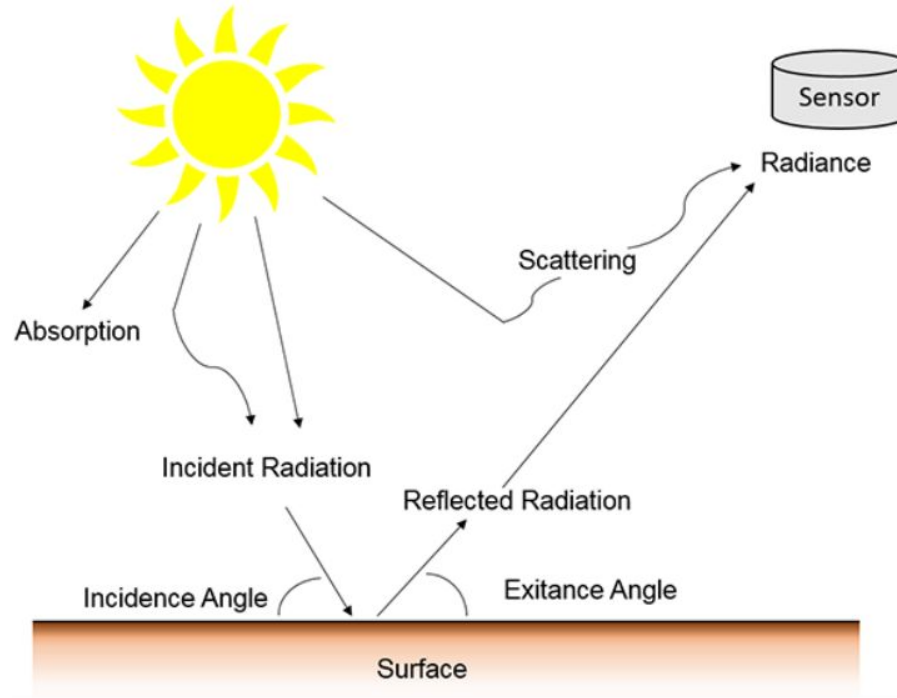


Image courtesy by Merav Mazouz & Matan Kolath

Reflectance model

In particular, we are interested in the light that is reflected by materials





Image courtesy by Merav Mazouz & Matan Kolath

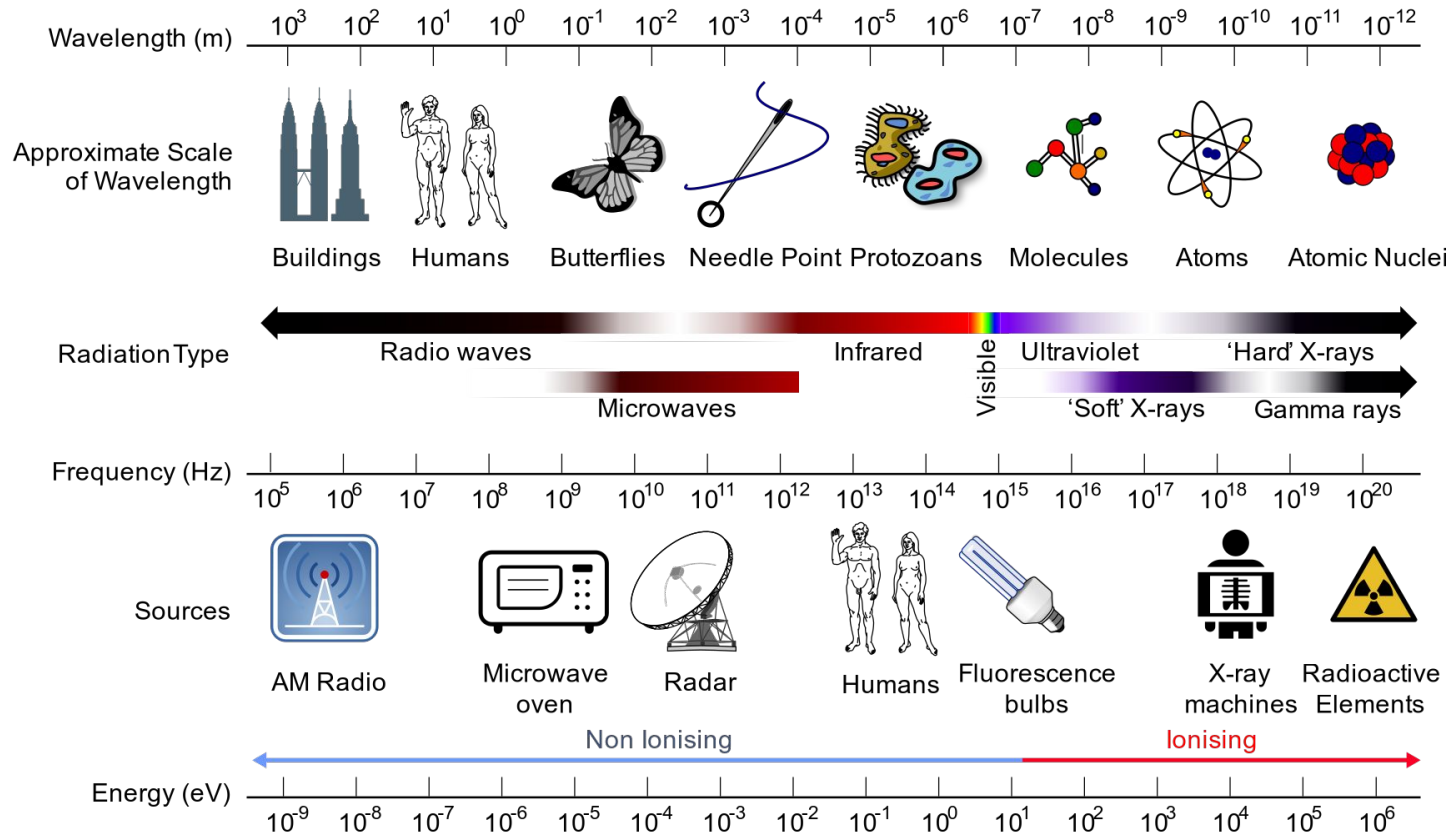
The electromagnetic spectrum

Types of Electromagnetic Radiation




© Encyclopædia Britannica, Inc.

The electromagnetic spectrum



The electromagnetic spectrum

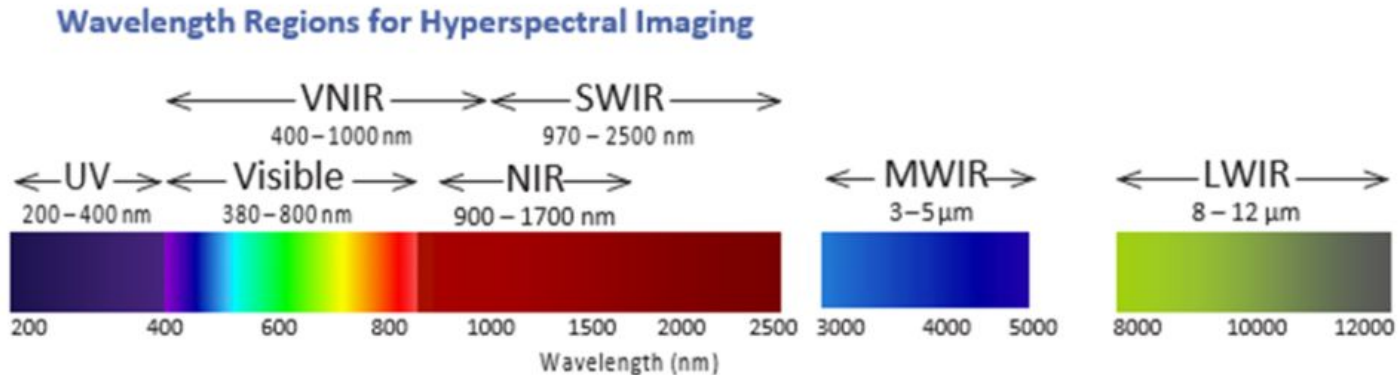
- Some animals can see a wider range of the spectrum



- Data doesn't exist solely in the visible light spectrum
- Hyperspectral imaging uses data from more than just the visible light spectrum



The hyperspectral spectrum

- It is composed of two portions of the EM spectrum:
 - **VNIR** = visible and near-infrared (VNIR) - [400, 1100] nm
 - **SWIR** = Short-wave infrared (SWIR) - [900 – 1700] nm
- VNIR includes the visible portion of the spectrum



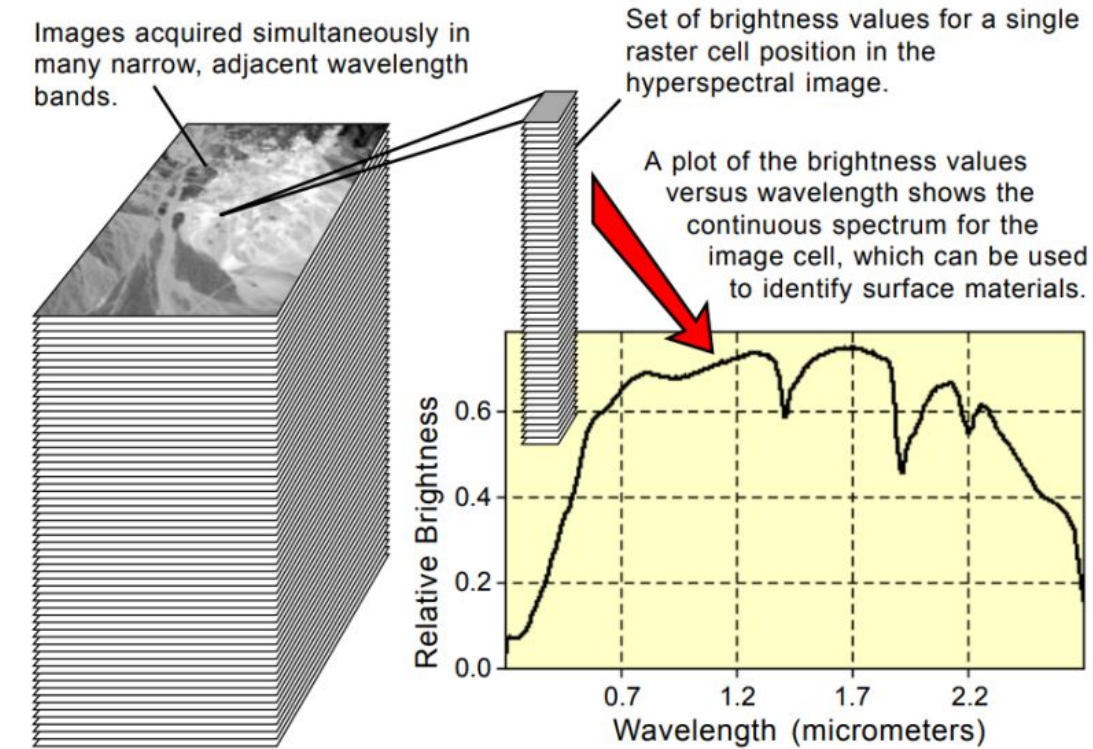
Spectral signature

- Every material has a unique spectral signature characterized by its reflectance graph



Hyperspectral imaging

Many narrow wavelength are measured in close proximity to create a continuous graph of wavelength to reflectance




Spectrograph



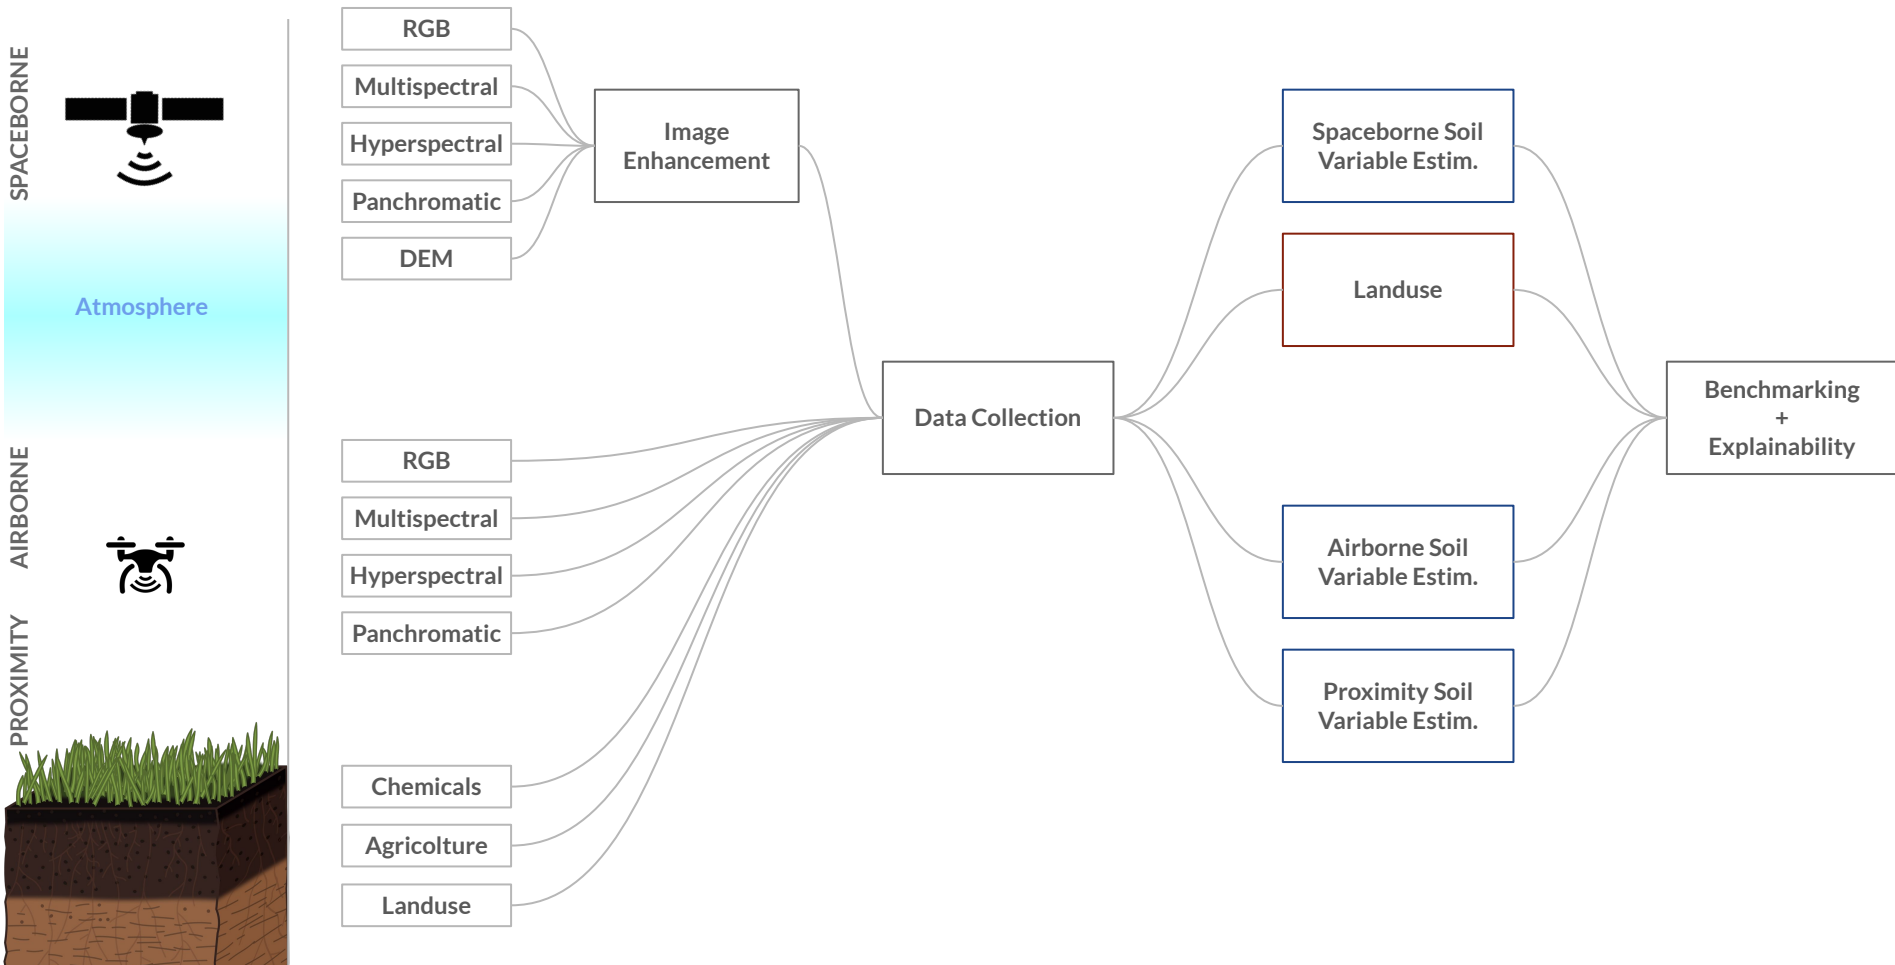
Digital soil mapping

- It is the computer-assisted production of digital maps of soil types and soil properties
- In the first place, several samples must be collected and analyzed in laboratory
 - it is very expensive
 - it is very time-consuming
- Then, it is possible to generate soil maps by interpolating on other locations



- Is it possible to create a deep-learning-based system to avoid manual analysis?
 - yes! Analyzing the hyperspectral signal
 - it is possible to mount the acquisition device on-board of drones and satellites

Setup



Lucas soil dataset


It's a dataset for digital soil mapping through the analysis of the hyperspectral signal

Each sample is composed of:

- **GPS position**
 - *GPS_LAT, GPS_LONG*
- **hyperspectral** signal 400 nm - 2499.5 nm
 - *spc.<WAVELENGTH>* columns
- **12 soil variables:**
 - *coarse, clay, silt, sand, pH.in.CaCl2, pH.in.H2O, OC, CaCO3, N, P, K, CEC*

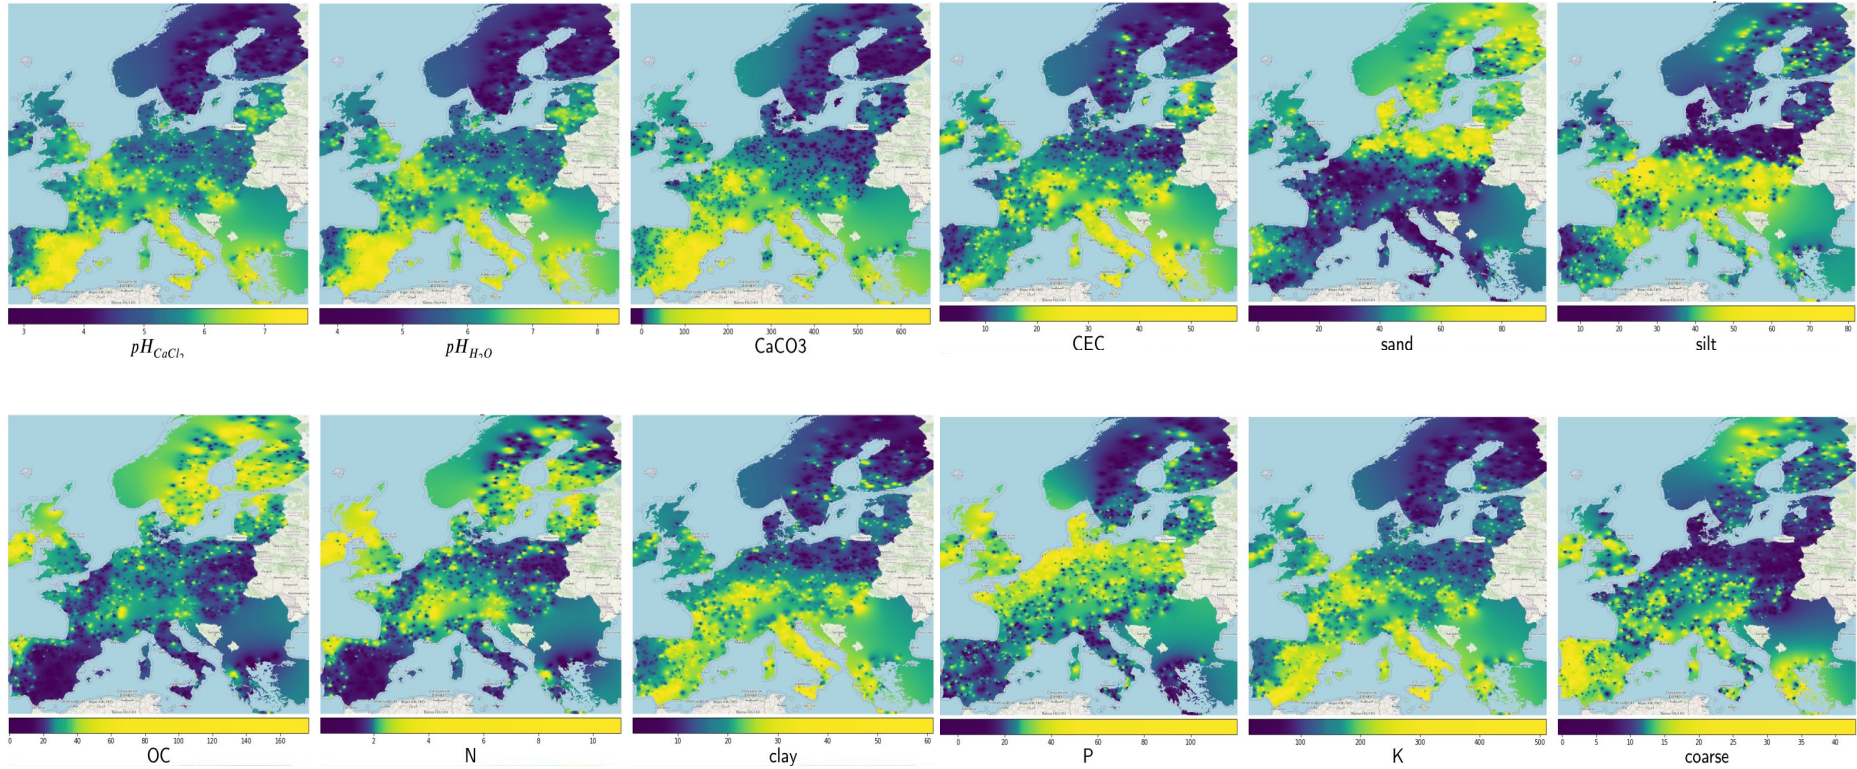
Lucas soil dataset

Example



Lucas soil dataset

Interpolation of the ground truths



Exercise

Train a convolutional neural network capable of estimating soil properties from hyperspectral signal

- On e-learning you will find:
 - the dataset
 - few lines of code needed to load the data in a proper way